

# Walmart commerce analysis and ARIMA forecasting

Brian Chen

Data from <https://www.kaggle.com/yassserh/walmart-dataset>.

Packages to be used.

```
library(tidyverse)
library(janitor)
library(lubridate)
library(scales)

options(scipen = 999)
```

Import data and do some light cleaning.

```
walmart = read.csv("walmart.csv")
head(walmart)

##   Store      Date Weekly_Sales
## 1     1 05-02-2010      1643691
## 2     1 12-02-2010      1641957
## 3     1 19-02-2010      1611968
## 4     1 26-02-2010      1409728
## 5     1 05-03-2010      1554807
## 6     1 12-03-2010      1439542
##   Holiday_Flag Temperature
## 1             0        42.31
## 2             1        38.51
## 3             0        39.93
## 4             0        46.63
## 5             0        46.50
## 6             0        57.79
##   Fuel_Price      CPI Unemployment
## 1     2.572 211.0964      8.106
## 2     2.548 211.2422      8.106
## 3     2.514 211.2891      8.106
## 4     2.561 211.3196      8.106
## 5     2.625 211.3501      8.106
## 6     2.667 211.3806      8.106

walmart$Date = dmy(walmart$Date)
walmart = clean_names(walmart)

head(walmart)
```

```

##   store      date weekly_sales
## 1     1 2010-02-05     1643691
## 2     1 2010-02-12     1641957
## 3     1 2010-02-19     1611968
## 4     1 2010-02-26     1409728
## 5     1 2010-03-05     1554807
## 6     1 2010-03-12     1439542
##   holiday_flag temperature
## 1             0        42.31
## 2             1        38.51
## 3             0        39.93
## 4             0        46.63
## 5             0        46.50
## 6             0        57.79
##   fuel_price      cpi unemployment
## 1     2.572 211.0964     8.106
## 2     2.548 211.2422     8.106
## 3     2.514 211.2891     8.106
## 4     2.561 211.3196     8.106
## 5     2.625 211.3501     8.106
## 6     2.667 211.3806     8.106

```

How many total weeks are being examined for each store?

```
walmart %>% count(store)
```

```

##   store n
## 1     1 143
## 2     2 143
## 3     3 143
## 4     4 143
## 5     5 143
## 6     6 143
## 7     7 143
## 8     8 143
## 9     9 143
## 10    10 143
## 11    11 143
## 12    12 143
## 13    13 143
## 14    14 143
## 15    15 143
## 16    16 143
## 17    17 143
## 18    18 143
## 19    19 143
## 20    20 143
## 21    21 143
## 22    22 143
## 23    23 143
## 24    24 143
## 25    25 143
## 26    26 143

```

```

## 27    27 143
## 28    28 143
## 29    29 143
## 30    30 143
## 31    31 143
## 32    32 143
## 33    33 143
## 34    34 143
## 35    35 143
## 36    36 143
## 37    37 143
## 38    38 143
## 39    39 143
## 40    40 143
## 41    41 143
## 42    42 143
## 43    43 143
## 44    44 143
## 45    45 143

```

Take a look at monthly sales for all years combined.

```

monthly_sales = walmart %>%
  group_by(store, month(date)) %>%
  summarize(total = sum(weekly_sales)) %>%
  rename(month = `month(date)`)
```

## ‘summarise()’ has grouped output by ‘store’. You can override using the ‘.groups’ argument.

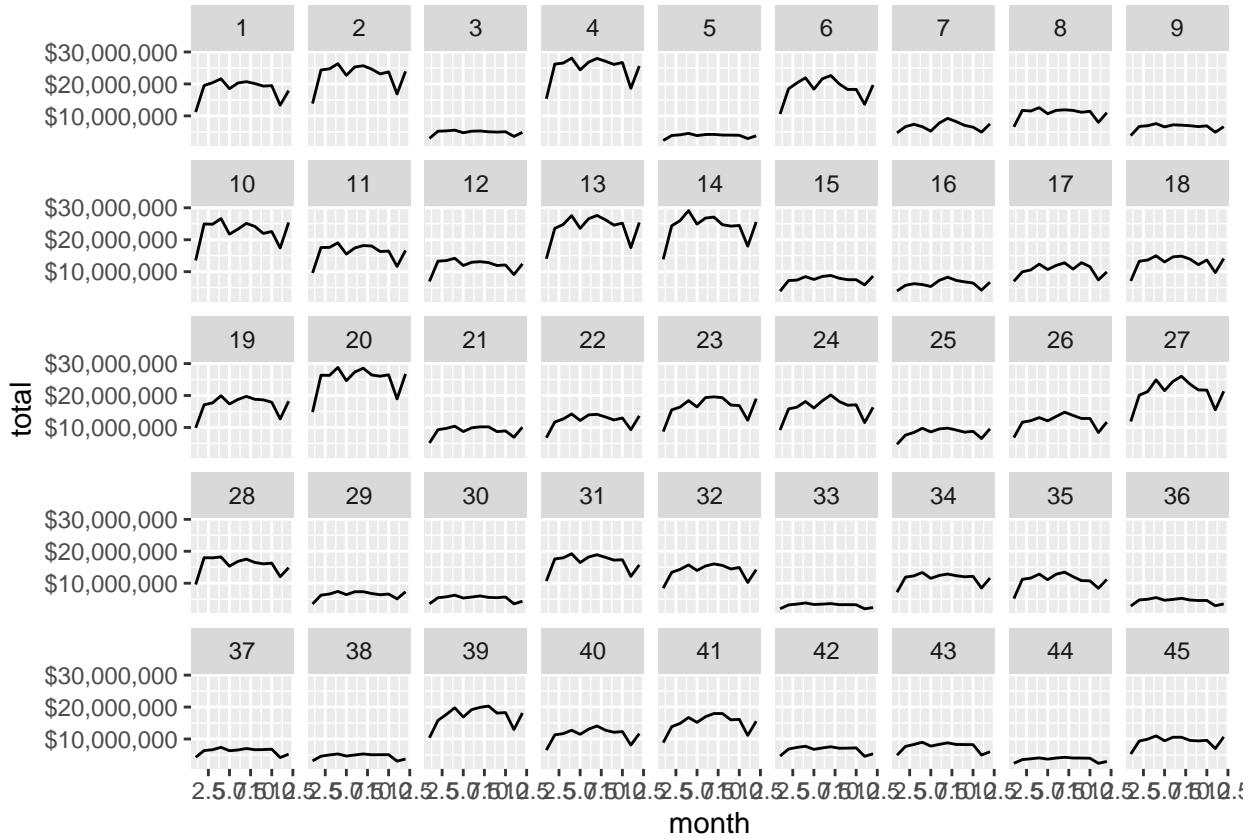
```
head(monthly_sales)
```

```

## # A tibble: 6 x 3
## # Groups:   store [1]
##   store month     total
##   <int> <dbl>     <dbl>
## 1     1     1 11203741.
## 2     1     2 19505307.
## 3     1     3 20380667.
## 4     1     4 21623140.
## 5     1     5 18505333.
## 6     1     6 20299636.
```

```

monthly_sales %>%
  ggplot(aes(x=month,y=total,group=store))+ 
  geom_line()+
  facet_wrap(~store, ncol=9)+ 
  scale_y_continuous(labels=scales::dollar_format())
```



Create function to show the most/least profitable store per year.

```
years = walmart %>%
  group_by(year(date)) %>%
  group_split()

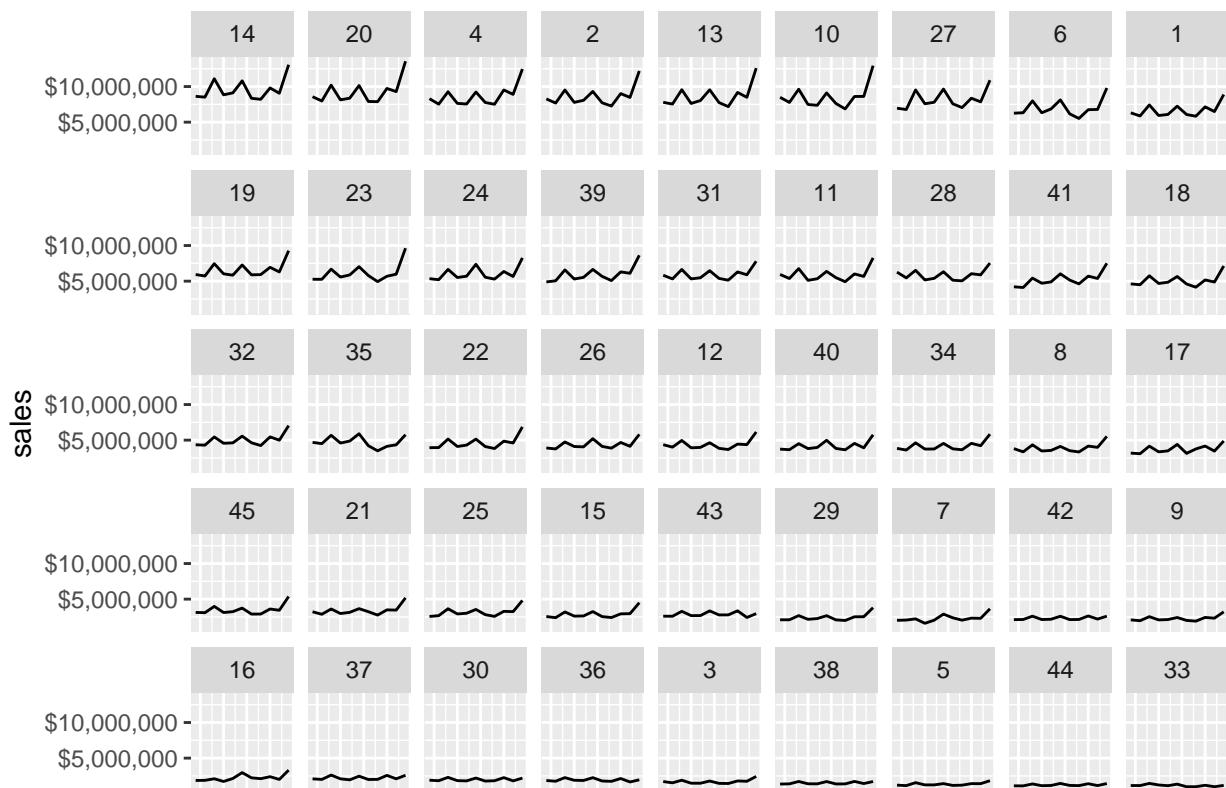
yearly_sales = function(year) {
  year %>%
  group_by(store, month(date)) %>%
  summarize(sales = sum(weekly_sales)) %>%
  ungroup() %>%
  group_by(store) %>%
  mutate(total = sum(sales)) %>%
  ungroup() %>%
  rename(month = `month(date)` %>%
  arrange(desc(total)) %>%
  ggplot(aes(x=month, y=sales, group=store))+
  geom_line()+
  facet_wrap(~reorder(store, -total), ncol=9)+
  scale_y_continuous(labels=scales::dollar_format())+
  theme(
    axis.text.x = element_blank(),
    axis.ticks.x = element_blank(),
    axis.title.x = element_blank()
  )+
  ggtitle(paste0("Most profitable stores in ",
                year$date[1] %>% str_extract("(201.)")))
```

```
}
```

```
yearly_sales(years[[1]]) #Caveat: data is missing January
```

```
## `summarise()` has grouped output by 'store'. You can override using the '.groups' argument.
```

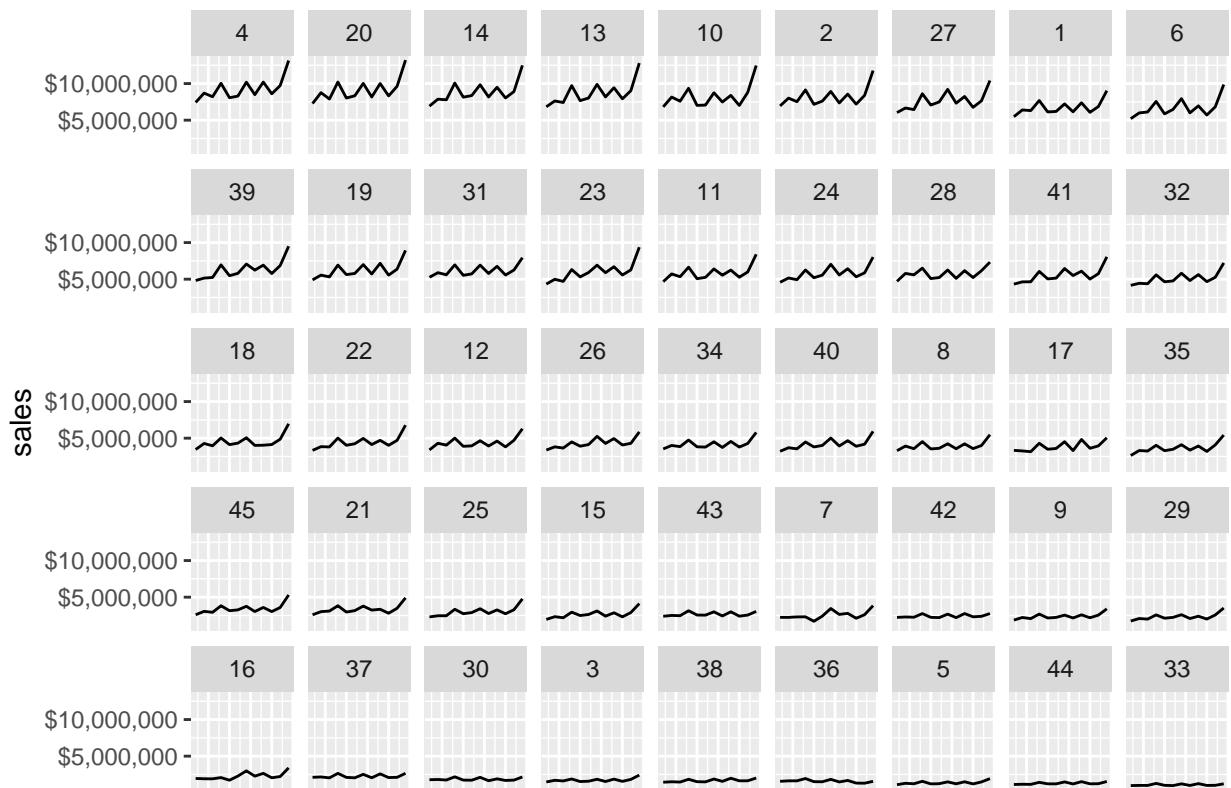
### Most profitable stores in 2010



```
yearly_sales(years[[2]])
```

```
## `summarise()` has grouped output by 'store'. You can override using the '.groups' argument.
```

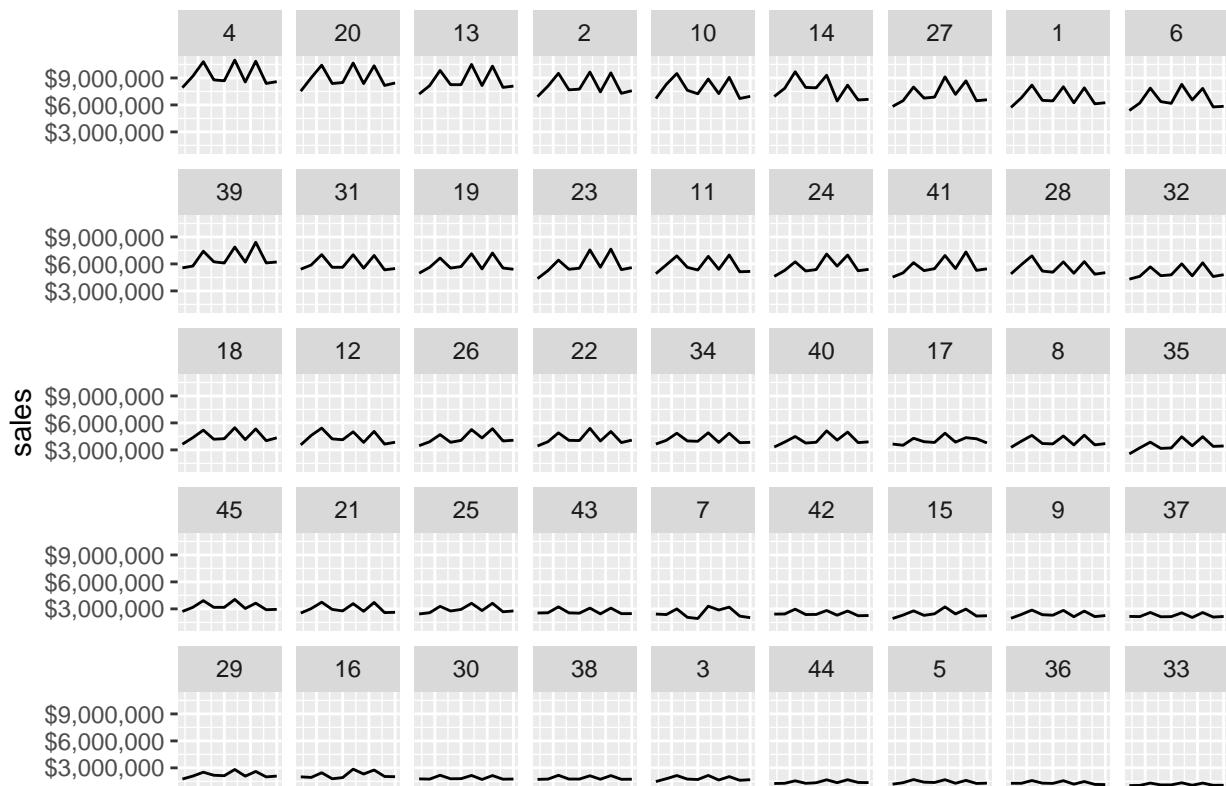
## Most profitable stores in 2011



```
yearly_sales(years[[3]]) #Caveat: data is missing November and December
```

```
## `summarise()` has grouped output by 'store'. You can override using the '.groups' argument.
```

## Most profitable stores in 2012



Checking overall performance of stores over the three years, check the ranking over each year, and check the best/worst performing stores over that period.

```
rankings = walmart %>%
  group_by(store, year(date)) %>%
  summarize(total = sum(weekly_sales)) %>%
  ungroup() %>%
  rename(year = `year(date)` %>%
  arrange(year, desc(total)) %>%
  mutate(score = rep(1:45,3)) #can assign score so easily b/c already arranged
```

## 'summarise()' has grouped output by 'store'. You can override using the '.groups' argument.

```
rankings %>%
  select(store, year, score) %>%
  pivot_wider(names_from = year, values_from = score)
```

```
## # A tibble: 45 x 4
##   store '2010' '2011' '2012'
##   <int>   <int>   <int>   <int>
## 1     1     14      1      3      6
## 2     2     20      2      2      2
## 3     3      4      3      1      1
## 4     4      2      4      6      4
## 5     5     13      5      4      3
```

```

##   6    10     6     5     5
##   7    27     7     7     7
##   8     6     8     9     9
##   9     1     9     8     8
## 10    19    10    11    12
## # ... with 35 more rows

rankings %>%
  group_by(store) %>%
  summarize(ranking = sum(score)) %>%
  arrange(ranking) #the lower the score, the better performing

## # A tibble: 45 x 2
##       store ranking
##       <int>    <int>
##   1      4        5
##   2     20       6
##   3     14      10
##   4     13      12
##   5      2      14
##   6     10      16
##   7     27      21
##   8      1      25
##   9      6      26
##  10    19      33
## # ... with 35 more rows

```

Checking the most profitable months.

```

#most profitable months?

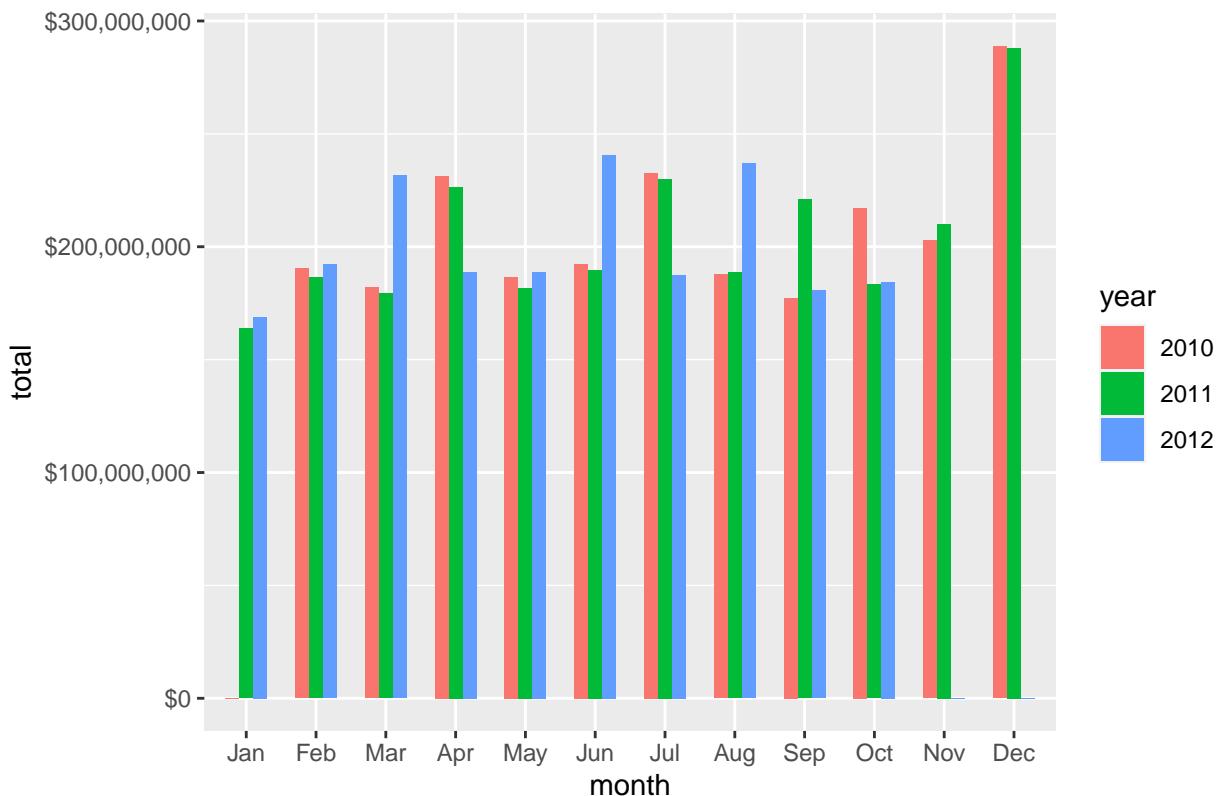
year_sales = walmart %>%
  group_by(year(date), month(date)) %>%
  summarize(total = sum(weekly_sales)) %>%
  rename(month = `month(date)`,
         year = `year(date)` %>%
  pivot_wider(names_from = year, values_from=total) %>% #doing this to use replace 0s
  arrange(month) %>%
  replace(is.na(.), 0) %>%
  pivot_longer(!month, names_to="year", values_to="total")

## 'summarise()' has grouped output by 'year(date)'. You can override using the '.groups' argument.

ggplot(year_sales, aes(x=month,y=total,fill=year))+ 
  geom_bar(position="dodge", width=0.6, stat="identity") +
  scale_x_discrete(limits = month.abb) +
  scale_y_continuous(labels=scales::dollar_format()) +
  ggtitle(label="Most profitable months, 2010-2012")

```

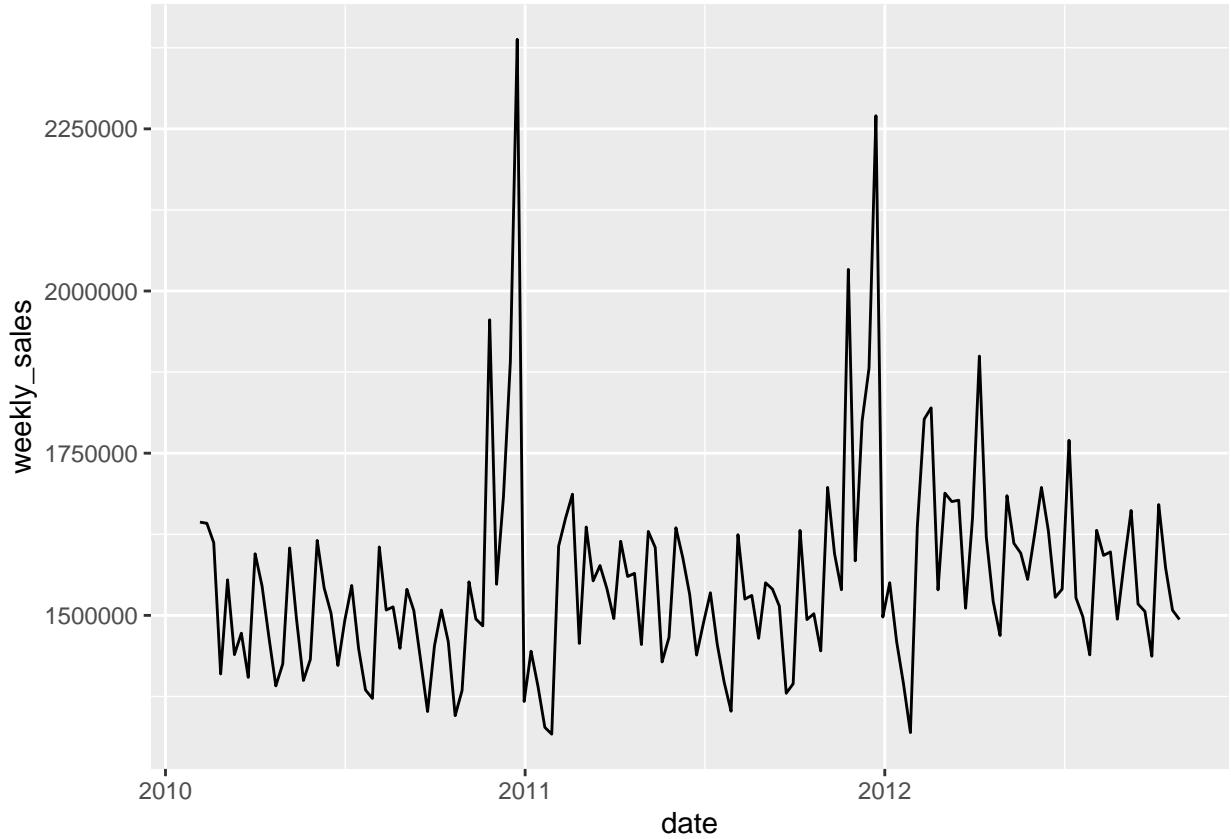
## Most profitable months, 2010–2012



Build ARIMA model for store 1

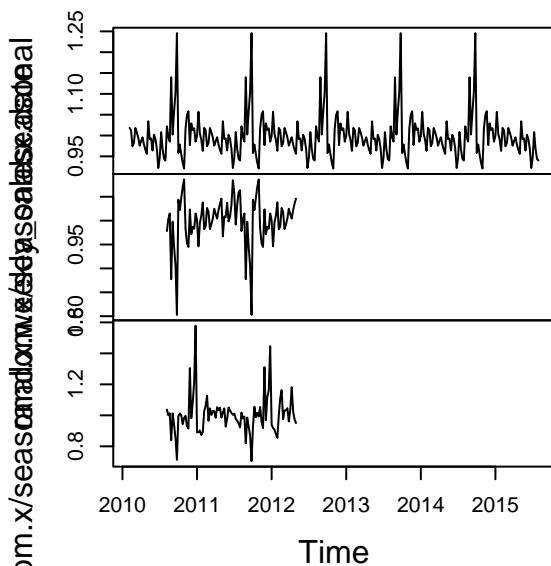
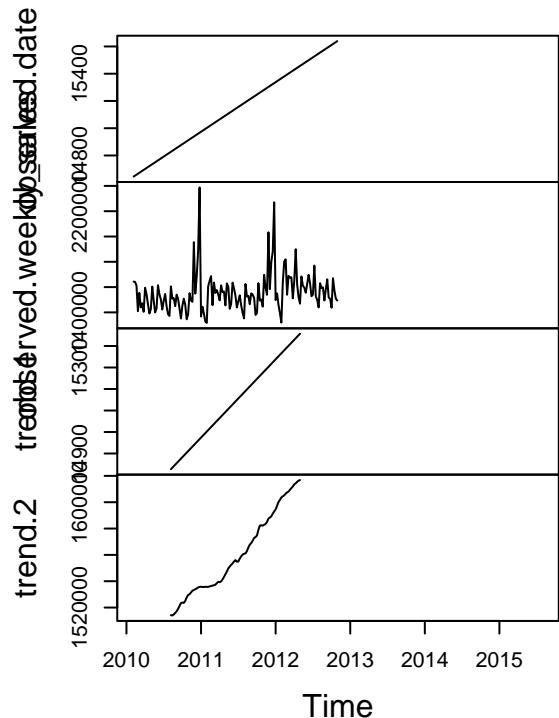
```
#Store 1 ARIMA

store1 = walmart %>% filter(store == 1) %>% select(date, weekly_sales)
ggplot(store1, aes(x=date, y=weekly_sales)) +
  geom_line()
```

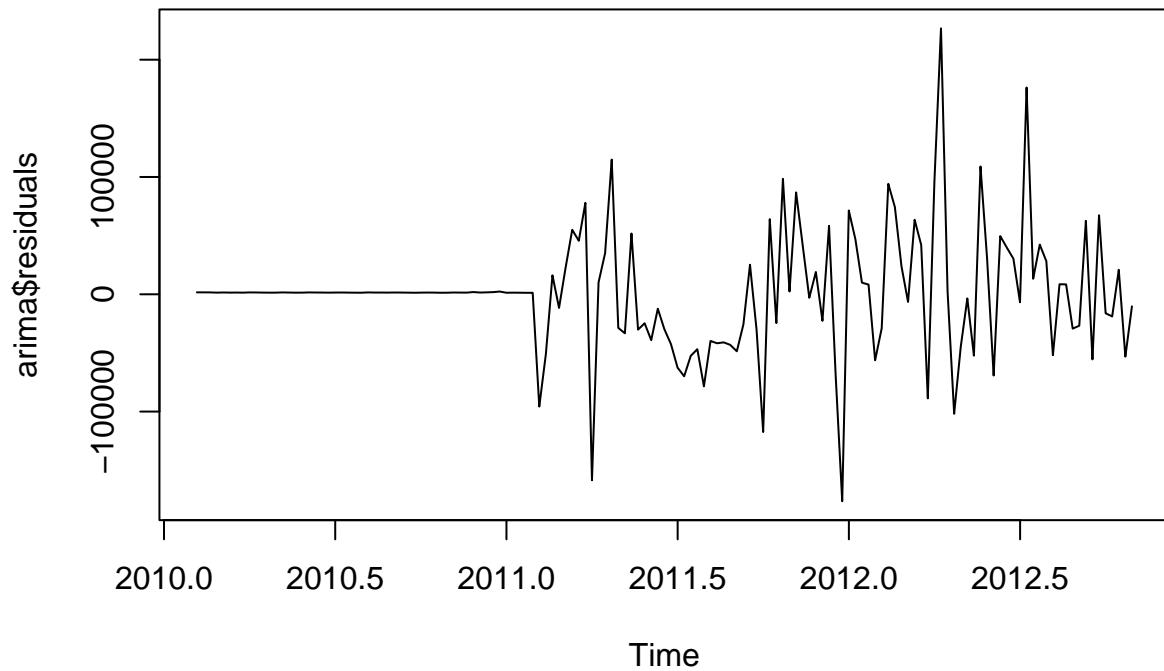


```
store1ts = ts(store1, frequency=52, start=c(decimal_date(ymd("2010-02-05"))))  
store1d = decompose(store1ts, "multiplicative")  
plot(store1d)
```

## Decomposition of multiplicative time series

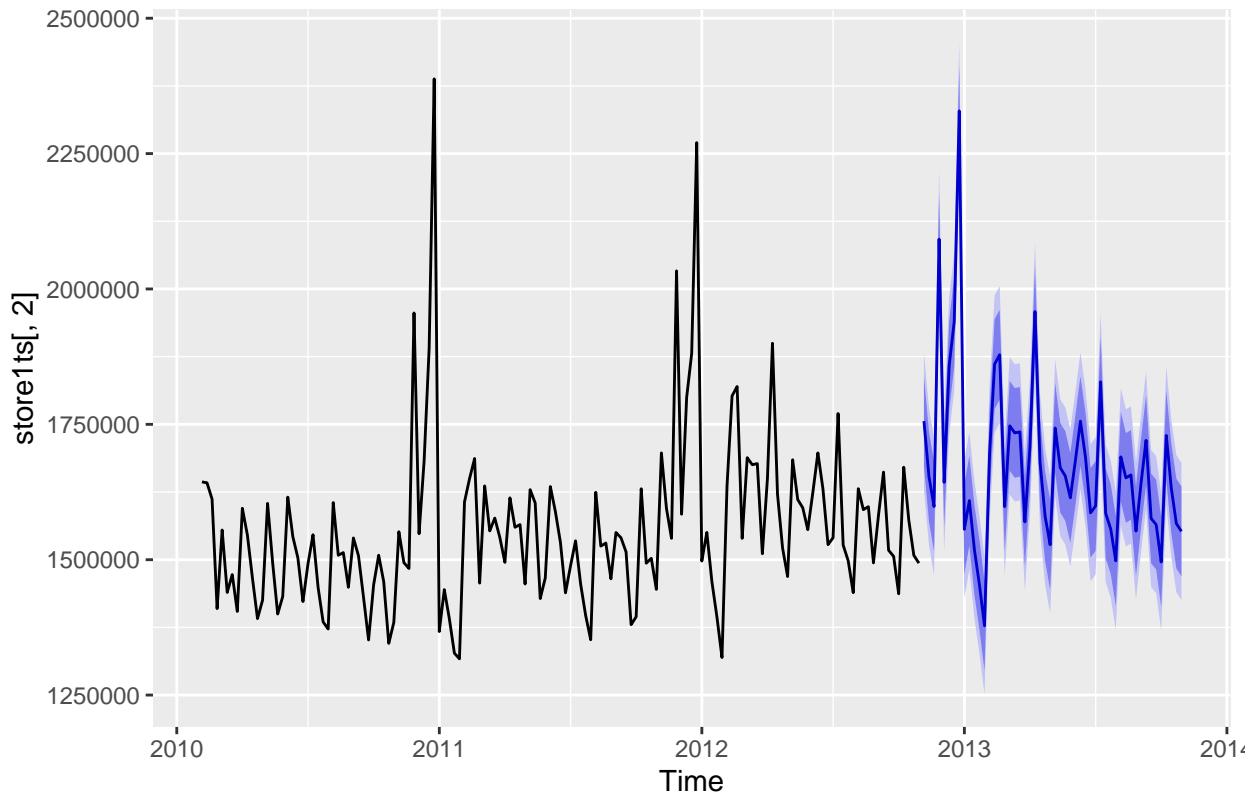


```
library(forecast)  
  
arima = auto.arima(store1ts[,2])  
plot.ts(arima$residuals)
```



```
forecast = forecast(arima, h=52)
autoplot(forecast)
```

### Forecasts from ARIMA(0,0,0)(0,1,0)[52] with drift



Using the Ljung-Box test to check accuracy.

```
Box.test(arima$resid, lag=5, type="Ljung-Box")
```

```
##
## Box-Ljung test
##
## data: arima$resid
## X-squared = 5.0543, df = 5,
## p-value = 0.4093
```

```
Box.test(arima$resid, lag=10, type="Ljung-Box")
```

```
##
## Box-Ljung test
##
## data: arima$resid
## X-squared = 9.6349, df = 10,
## p-value = 0.4731
```

```
Box.test(arima$resid, lag=15, type="Ljung-Box")
```

```
##
## Box-Ljung test
```

```

##  

## data: arima$resid  

## X-squared = 22.434, df = 15,  

## p-value = 0.09693

```

Writing a reusable ARIMA function.

```

store_arima = function(storeNum, weeks_forecasted) {  

  df = walmart %>%  

    filter(store == storeNum) %>%  

    select(date, weekly_sales)  

  print(ggplot(df, aes(x=date, y=weekly_sales)) +  

    geom_line() +  

    ggtitle(label=paste0("Sales for Store ", storeNum)))  

  df_ts = ts(df, frequency=52, start=c(decimal_date(ymd("2010-02-05"))))  

  df_d = decompose(df_ts, "multiplicative") %>% plot()  

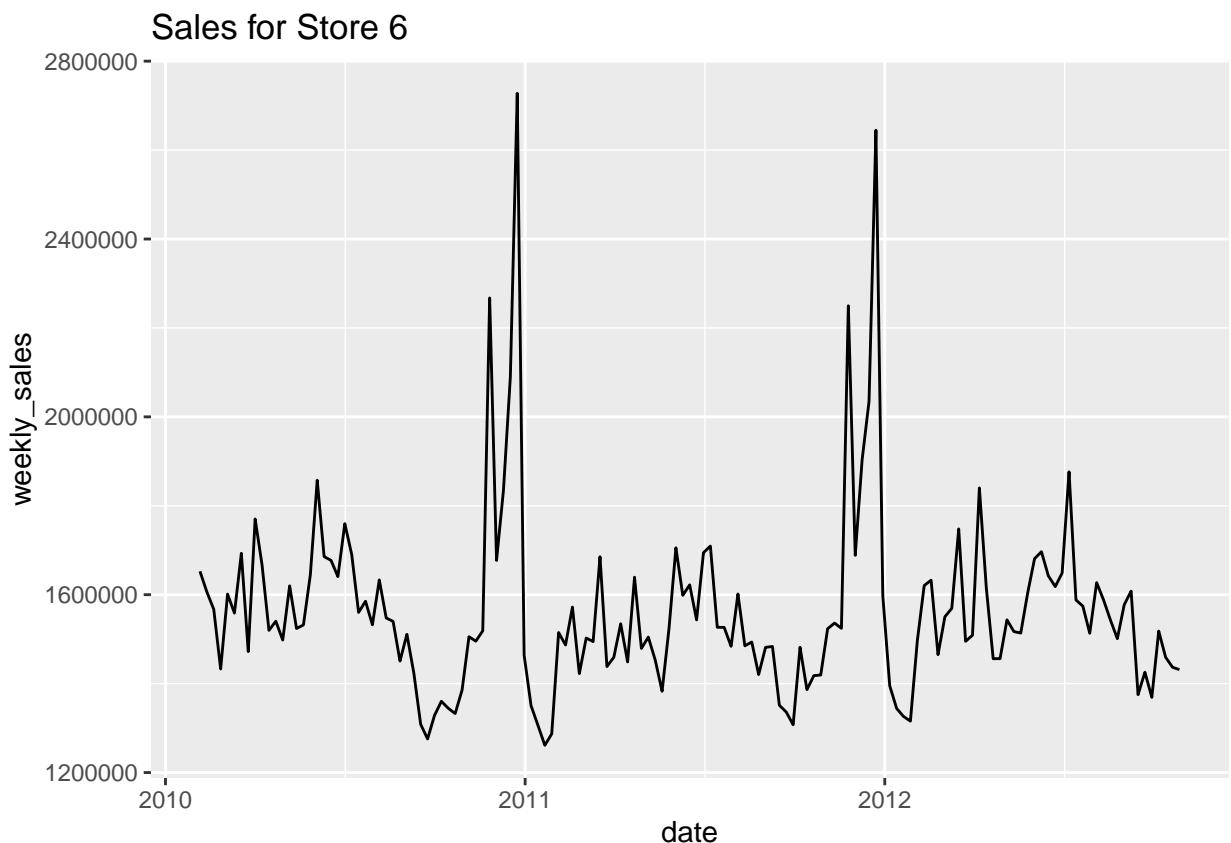
  forecast(auto.arima(df_ts[,2]), h=weeks_forecasted) %>% autoplot()  

}  

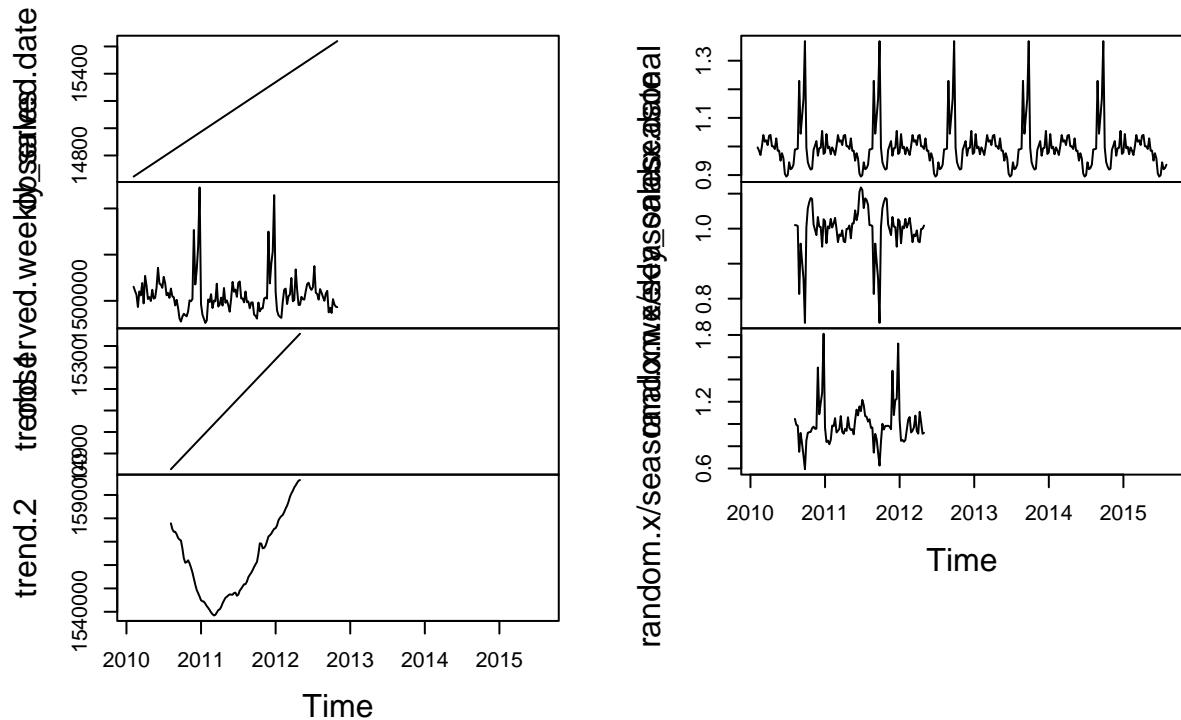
  

store_arima(6,20) #store_arima(which store, forecast how many weeks)

```



## Decomposition of multiplicative time series



Forecasts from ARIMA(1,1,3)(0,1,0)[52]

